

## Operating System HW1

**(1) On all modern computers, at least part of the interrupt handlers are written in assembly language. Why?**

Because when an interrupt occur, there are corresponding actions have to happen. Those actions cannot be written by high level languages.

**(2) If a multi-threaded process forks, a problem occurs if the child gets copies of all the parents threads. Suppose that one of the original threads was waiting for keyboard input. Now two threads are writing for keyboard input, one in each process. Does this problem ever occur in single-threaded processes?**

No, the read() call will make the calling thread block until the keyboard inputs arrive, in the multithreaded case, one thread in parent process is blocked because of read() call, then another thread call fork(), the child process will be created and copy the parent states, which means the read() call will also be copied. when the keyboard inputs arrive, from the view of OS, only one thread will be unblocked, either read() in parent or child.

while in single thread case, there is only one execution unit, no matter at what time there are only one read() call, that's why the single-threaded process will not have such problem.

**(3) Why would a user thread ever voluntarily give up the CPU by calling thread yield? After all, since there is no periodic clock interrupt, it may never get the CPU back.**

Because it is the user programmer who writes the code that manage the user threads in one process, there is no thread "scheduler" for those threads. To prevent a user thread take too much CPU time, it is best for the programmer to do such job.

**(4) Consider a process that continuously reads in a chunk of data from disk (1MB say), does a heavy computation over that 1MB and then reads in the next chunk (adinfinitum). How would this Multi-Level Feedback queues handle this process?**

Since this process will take a lot of CPU time and is also I/O-bound, when it is dealing with I/O-bound work, it will be moved to higher priority queues, and when it do the computation work, it will be moved to lower-priority queue. However, if it stays in a lower-priority queue for a long time, it may be moved to a higher priority-queue.

**(5)**

- (a) kernel-level threading, because in this case, it is the OS managing the threads, All thread operations are implemented in the kernel and the OS schedules all threads in the system
- (b) no-threading. Because threading will not help much in this case.
- (c) hybrid threading, for it can combine the advantages of user-thread and kernel thread to satisfy the requirements of this application.

**(6) Discuss why the solution in which every philosopher takes and locks the fork on there left and then tries to grab the fork on their right fails.**

Suppose every philosopher takes and locks the fork on the left at the same time, then when they tries to grab the fork on the right, they will find there is no fork available, deadlock occurs.

**(7)**

**a) Describe the purpose for each of the global variables below:**

i) reindeerSem:

Controlling the synchronization and mutex between Santa() and reindeer(). As we can see in the code, at the beginning, each reindeer will execute reindeer.P() and then get blocked, and when the Santa() wake up, the Santa() will call reindeer.V() to wake up reindeer, later the getHiched() is able to execute.

ii) santaSem:

Controlling the synchronization and mutex between Santa() and reindeer(). At the beginning, the Santa() will get blocked because of santaSem.P(). When the reindeer number equals to 9, then the reindeer() will call santaSem.V() to wake the Santa up.

iii) mutex:

“mutex” ensure only one process can enter the critical section at each time. This variable initialized as 1, because at first there must be one process is able to get into the critical section.

iv) reindeer:

This variable indicated the number of reindeers. Each reindeer process will revise this variable when it get into the critical section, it this variable equals 9 which means the ninth reindeer arrives.

**b) Write the elf portion of the synchronization routine.**

// there are N elves.

```
void elf(){
    elfTex.P();
    mutex.P();
```

```

elves+=1;
if(elves == 3){
    santaSem.V();
}
else{
    elfTex.V();
}
mutex.V();
getHelp();
mutex.P();
elves -=1;
if(elves == 0) elfTex.V();
mutex.V();
}

```

**(8) Given a 48-bit addressable CPU, a page size of 8KB, a page entry size of 8 bytes, and a ram size of 4GB answer the following questions. You can leave your answers in the form  $2^x$  (where you would give x).**

**a) How big is the (maximum) virtual address space?**

Because we are given the 48-bit addressable CPU, the virtual address space will be  $2^{48}$  bytes.

**b) How many entries are there in the page table (assuming the maximum virtual address space size)?**

Page Size = 8KB =  $2^{13}$ B

Offset = 13 bit

number of page entries =  $2^{(48-13)} = 2^{35}$

**c) How large (in bytes) would the page table be?**

page table size = number of page entries \* page entry size =  $2^{35} * 8(\text{bytes}) = 2^{38}\text{byte}$

**d) How large (in bytes) would an inverted page table be?**

number of pages =  $2^{48}/2^{13} = 2^{35}$

number of bits required to identify each page = 35 bits (5 Bytes)

number of frames =  $2^{32}/2^{13} = 2^{19}$

inverted page table size = number of frames \* number of bits required to identify each page =  $2^{19} * 5 \text{ Bytes} = 2560 \text{ KB}$

**e) Using a two level multi-level page table where each second- level page table would be a page-size large, what would an address look like? (i.e. How many bits for offset, 2nd level and top level fields?)**

13 bits for offset because page size is still 8KB.

since the 2nd level page table size would be a page-size large which equals to 8KB, there will be 13 bits for 2nd level page table. we can calculate 2nd level offset by:  
 (page size/ page entry size) = 8KB / 8B =  $2^{10}$ , there will be 10 bits for offset.

Then we have  $48 - 13 - 10 = 25$  bits to address top level field.

**f) Explain how translation from a virtual address to a physical address occurs using a multi-level page table.**

Take a two-level page table as an example:

the virtual address can be expressed as:

|      P1      |      P2      |      offset      |

The P1 is used to index the 2nd level page tables. From P1, we can get which one among 2nd level page tables we should look at.

The P2 is the offset of the 2nd level page table. From P2, we can get which page we should look at.

The offset is used to get the which PTE we should access in one page table;

Then we can get physical frame number through the PTE, and the physical address equals to: physical frame number with offset.

**(9)**

(a) FIFO

sequence	Frame 1	Frame 2	Frame 3	Frame 4	Page Fault
read P0	P0				Y
write P2	P0	P2			Y
read P3	P0	P2	P3		Y
write P4	P0	P2	P3	P4	Y
read P5	P5	P2	P3	P4	Y
write P3	P5	P2	P3	P4	N
read P0	P5	P0	P3	P4	Y
write P2	P5	P0	P2	P4	Y

Number of page fault: 7

(b) optimize:

sequence	Frame 1	Frame 2	Frame 3	Frame 4	Page Fault
read P0	P0				Y
write P2	P0	P2			Y
read P3	P0	P2	P3		Y
write P4	P0	P2	P3	P4	Y
read P5	P0	P2	P3	P5	Y
write P3	P0	P2	P3	P5	N
read P0	P0	P2	P3	P5	N
write P2	P0	P2	P3	P5	N

Number of page fault: 5

(c) LRU

sequence	Frame 1	Frame 2	Frame 3	Frame 4	Page Fault
read P0	P0				Y
write P2	P0	P2			Y
read P3	P0	P2	P3		Y
write P4	P0	P2	P3	P4	Y
read P5	P5	P2	P3	P4	Y
write P3	P5	P2	P3	P4	N
read P0	P5	P0	P3	P4	Y
write P2	P5	P0	P3	P2	Y

Number of page fault: 7

(10)

**a, How many addresses fit in an i-node?**

First, there are 5 direct addresses.

To calculate single indirect: we have 4 bytes per disk address, and the block size is 32 bytes, then the single direct contain  $32\text{bytes}/4\text{bytes} = 8$  blocks

For the double indirect, it will contain  $8 \times 8 = 64$  blocks

In conclusion, there are  $5 + 8 + 64 = 77$  addresses fit in an i-node.

**b, What is the maximum file-size?**

direct:  $5 \times 32 \text{ bytes} = 160 \text{ bytes}$

$32\text{bytes}/4\text{bytes} = 8 \text{ blocks}$

single indirects:  $8 \times 32\text{bytes} = 256 \text{ bytes}$

double indirects:  $8 \times 8 \times 32\text{bytes} = 2048 \text{ bytes}$

maximum file size:  $160 + 256 + 2048 = 2464 \text{ bytes}$

**c, How many (maximally-sized) files can fit on the disk?**

disk address = 4 bytes = 32 bits;

the disk space =  $2^{32} = 4\text{GB}$

thus there are  $4\text{GB}/2464\text{B} = 1743087$  (maximally-sized)files

**d, Use the example in Figure 2 to convert the following file offsets (in bytes) to physical block addresses:**

**100, 301, 633**

100:  $100/32 = 3 \dots 4$ , so the physical block address is saved in block#4, and thus equals to 13.

301:  $301/32 = 9 \dots 13$ , so the physical block address is saved in block#10, and thus equals to 228.

633:  $633/32 = 19 \dots 25$ , so the physical block address is saved in block#20, and thus equals to 8621.